

Data manipulation with DataFrames

INTRODUCTION TO PYSPARK



Ben Schmidt
Data Engineer

Handling missing data

- Use `.na.drop()` to remove rows with null values

```
# Drop rows with any nulls
```

```
df_cleaned = df.na.drop()
```

```
# Filter out nulls
```

```
df_cleaned = df.where(col("columnName").isNotNull())
```

- Use `.na.fill({"column": value})` to replace nulls with a specific value

```
# Fill nulls in the age column with the value 0
```

```
df_filled = df.na.fill({"age": 0})
```

Column operations

- Use `.withColumn()` to add a new column based on calculations or existing columns

```
# Create a new column 'age_plus_5'  
df = df.withColumn("age_plus_5", df["age"] + 5)
```

- Use `withColumnRenamed()` to rename columns

```
# Rename the 'age' column to 'years'  
df = df.withColumnRenamed("age", "years")
```

- Use `drop()` to remove unnecessary columns

```
# Drop the 'department' column  
df = df.drop("department")
```

Row operations

- Use `.filter()` to select rows based on specific conditions

```
# Filter rows where salary is greater than 50000
filtered_df = df.filter(df["salary"] > 50000)
```

- Use `.groupBy()` and aggregate functions (e.g., `.sum()`, `.avg()`) to summarize data

```
# Group by department and calculate the average salary
grouped_df = df.groupBy("department").avg("salary")
```

Row Operations Outcomes

- Filtering

```
+-----+----+-----+
|salary|age|      occupation  |
+-----+----+-----+
| 60000| 45|Exec-managerial  |
| 70000| 35|Prof-specialty   |
+-----+----+-----+
```

- GroupBy ```

```
+-----+-----+ |department|avg(salary)| +-----+-----+ | HR| 80000.0| |
IT| 70000.0| +-----+-----+
```

```

# CheatSheet

```
Drop rows with any nulls
df_cleaned = df.na.drop()

Drop nulls on a column
df_cleaned = df.where(col("columnName").isNotNull())

Fill nulls in the age column with the value 0
df_filled = df.na.fill({"age": 0})
```

- Use `.withColumn()` to add a new column based on calculations or existing columns.  
Syntax: `.withColumn("new_col_name", "original transformation")`

```
Create a new column 'age_plus_5'
df = df.withColumn("age_plus_5", df["age"] + 5)
```

- Use `withColumnRenamed()` to rename columns Syntax: `withColumnRenamed( old column name , new column name``

```
Rename the 'age' column to 'years'
df = df.withColumnRenamed("age", "years")
```

- Use `drop()` to remove unnecessary columns Syntax: `.drop( column name )`

# Let's practice!

INTRODUCTION TO PYSPARK

# Advanced DataFrame operations

INTRODUCTION TO PYSPARK



**Ben Schmidt**  
Data Engineer



# Joins in PySpark

- Combine rows from two or more DataFrames based on common columns
- Types of joins: inner, left, right, and outer, like SQL
- Syntax: `DataFrame1.join(DataFrame2, on="column", how="join_type")`

```
Joining on id column using an inner join
df_joined = df1.join(df2, on="id", how="inner")

Joining on columns with different names
df_joined = df1.join(df2,
 df1.Id == df2.Name,
 "inner")
```

# Union operation

- Combines rows from two DataFrames with the same schema
- Syntax: `DataFrame1.union(DataFrame2)`

```
Union of two DataFrames with identical schemas
df_union = df1.union(df2)
```

# Working with Arrays and Maps

**Arrays:** Useful for storing lists within columns, syntax: `ArrayType(StringType(), False)`

```
from pyspark.sql.functions import array, struct, lit

Create an array column
df = df.withColumn("scores", array(lit(85), lit(90), lit(78)))
```

**Maps:** Key-value pairs, helpful for dictionary-like data, `MapType(StringType(), StringType())`

```
from pyspark.sql.types import StructField, StructType, StringType, MapType

schema = StructType([
 StructField('name', StringType(), True),
 StructField('properties', MapType(StringType(), StringType()), True)
])
```

# Working with Structs

- Structs: Create nested structures within rows Syntax:

```
StructType(Structfield, Datatype())
```

```
Create a struct column
```

```
df = df.withColumn("name_struct", struct("first_name", "last_name"))
```

```
Create a struct column
```

```
df = df.withColumn("name_struct", struct("first_name", "last_name"))
```

# Let's practice!

INTRODUCTION TO PYSPARK

# U define it? U use it!

INTRODUCTION TO PYSPARK



**Benjamin Schmidt**  
Data Engineer

# UDFs for repeatable tasks

**UDF (User-Defined Function):** custom function to work with data using PySpark dataframes

Advantages of UDFs:

- Reuse and repeat common tasks
- Registered directly with Spark and can be shared
- PySpark DataFrames (for smaller datasets)
- pandas UDFs (for larger datasets)

# Defining and registering a UDF

All PySpark UDFs need to be registered via the `udf()` function.

```
Define the function
def to_uppercase(s):
 return s.upper() if s else None

Register the function
to_uppercase_udf = udf(to_uppercase, StringType())

Apply the UDF to the DataFrame
df = df.withColumn("name_upper", to_uppercase_udf(df["name"]))

See the results
df.show()
```

**Remember:** UDFs allow you to apply custom Python logic on PySpark DataFrames



# pandas UDF

- Eliminates costly conversions of code and data
- Does not need to be registered to the SparkSession
- Uses pandas capabilities on extremely large datasets

```
from pyspark.sql.functions import pandas_udf
```

```
@pandas_udf("float")
```

```
def fahrenheit_to_celsius_pandas(temp_f):
```

```
 return (temp_f - 32) * 5.0/9.0
```

# PySpark UDFS vs. pandas UDFs

## PySpark UDF

- Best for relatively small datasets
- Simple transformations like data cleaning
- Changes occur at the columnar level, not the row level
- Must be registered to a Spark Session with `udf()`

## pandas UDF

- Relatively large datasets
- Complex operations beyond simple data cleaning
- Specific row level changes over column level
- Can be called outside the Spark Session

# Let's practice!

INTRODUCTION TO PYSPARK